

COMPSCI 230 Assignment 2 (2019 S2)

Learning outcomes of this assignment

In completing this assignment, you will gain experience in:

- Designing the class structure for a Java GUI application with Swing components (a menu, radio buttons, a combo box, a JPanel, JTable and a custom table model – a subclass of AbstractTableModel)
- Using a JFileChooser
- Handling various events associated with Swing components
- Data input from a text file
- Use event-dispatch thread for GUI multi-threading

Depending on your design, you may also be using:

- Regular expressions
- Polymorphism with an abstract class and other OO features
- Exposure to Set collections (not taught in lectures)

This maps to the following learning outcomes of COMPSCI 230:

- *OO Programming*: describe and use the features typically offered by an object-oriented programming language, including support for classes, visibility, inheritance, interfaces, polymorphism and dynamic binding
- *OO Design*: explain and apply key design principles of object-oriented software development, including separation of concerns, abstraction, information hiding, programming to interfaces, coupling and cohesion, resilience to change, and reuse
- create simple OO design models
- *Frameworks*: describe important concepts of programming frameworks, including APIs, inversion of control and event-driven programming
- use a framework to develop a multithreaded GUI application
- *Concurrent Programming*: explain and apply the principles of application-level multithreading: threading, condition synchronization, mutual exclusion, and primitives associated with these

This assignment will be worth **5% of your final mark**.

Note: You can attempt up to 120 marks in this assignment, but you will only require any 100 of these to obtain full marks (i.e., there are 20 bonus marks that allow you to make up for marks lost in other parts of the assignment).

Due date and submission

Due: **11:59 pm on Oct 25, 2019**

Submission: **via the assignment dropbox** <https://adb.auckland.ac.nz>

What to submit: **all .java files that make up your application.** Do not submit as a ZIP file.

Introduction

In this assignment, you will work on a real life research problem. Our network lab hosts the Auckland Satellite Simulator (<http://sde.blogs.auckland.ac.nz/>). We use it to simulate Internet traffic to small Pacific islands that are connected to the rest of the world via a satellite link.

To do this, the simulator has a number of machines ("source hosts") on the "world side" of the simulated satellite link that transmit data in the form of small chunks of up to 1500 bytes called *packets*. These packets travel via a simulated satellite link to the "island side". Once on the

island side, each packet ends up at a machine there. These machines are called “destination hosts”.

The simulated satellite link delays packets and occasionally throws some away when there are more packets arriving that it can deal with at the moment. When the link throws packets away, the source hosts respond by sending less data for a while before attempting to send more again.

On the island side of the satellite link, our simulator eavesdrops on the incoming packets. It stores summary information about each packet in a *trace file*. The trace file is plain text and each line contains the record for exactly one packet (more on the file format in the next section).

What we would like to be able to do is get a tabular display of how much data comes from a particular source host over time, or how much data goes to a particular destination host over the course of an experiment. Experiments typically take between 90 seconds and about 11 minutes.

This is where your assignment comes in: You are to build an application that displays this data in tabular form.

In the next section, we’ll look at the structure of our trace files, before we take you through the steps for building the application.

Trace files

The total size of a trace file can vary substantially depending on the number of hosts involved in an experiment and the length of the experiment. On Canvas, there are two trace files for you to experiment with: a small one with 148315 lines, and a large one with 651274 lines. You can open them in text editors such as Notepad++. And yes you can edit them, too!

To get a good idea of your “typical” line, scroll down a good bit – the lines/packets at the start and at the end are a bit unusual. The following shows a bunch of typical lines from the large trace file:

108860	128.879102000	192.168.0.24	47928	10.0.0.5	5201	1514	1500	1448	...
108861	128.879885000	192.168.0.24	47928	10.0.0.5	5201	1514	1500	1448	...
108862	128.880603000	192.168.0.24	47928	10.0.0.5	5201	1514	1500	1448	...
108863	128.881481000	192.168.0.15	8000	10.0.1.25	59590	66	52	0	...
108864	128.881481000	192.168.0.9	8000	10.0.1.15	42081	66	52	0	...
108865	128.881481000	192.168.0.24	47928	10.0.0.5	5201	1514	1500	1448	...
108866	128.881495000	192.168.0.15	8000	10.0.1.25	59590	66	52	0	...
108867	128.882148000	192.168.0.3	8000	10.0.1.4	55442	1514	1500	1448	...
108868	128.882905000	192.168.0.3	8000	10.0.1.4	55442	1514	1500	1448	...
108869	128.883800000	192.168.0.3	8000	10.0.1.4	55442	1514	1500	1448	...

The lines here have been truncated to be able to accommodate them on the page, but they show all the data you will need. Each line consists of a number of fields separated by a single tab character. Note that fields may be empty, in which case you get two successive tab characters.

The **first field** on the left is just a sequential number for each packet that is added by the eavesdropping program. The **second field** is a time stamp that is also added by the eavesdropping program. Each trace file starts with a time stamp of 0.000000000 for the first packet in the file.

The **third field** in each line is the IP address of the source host. IP addresses identify machines on the network and help routers forward packets between source and destination. Each IP address consists of four decimal numbers between 0 and 255 separated by dots (full stops).

The trace files here only show packets heading towards destination hosts on the island side, so all source host IP addresses start with "192.168.0.", indicating that they are "world side" addresses.

The **fifth field** is the IP address of the destination host from the island network. All of the island addresses start with "10.0.". You will need the third or the fifth field to populate the combo box depending on the status of the radio buttons.

The **fourth** and the **sixth field** are the TCP ports on the hosts that the respective packets travel between. They identify the applications that have sent or would have received the packets, but are not relevant for your assignment.

Fields **seven, eight** and **nine** are packet sizes in bytes. The size we're interested in here is that in **field eight**, it's the IP packet size. The size in **field seven** is that of the whole Ethernet frame that contains the IP packet, and **field nine** is the TCP payload size (the size of the content of the IP packet).

There are also a number of **additional fields**: various flags, packet sequence and acknowledgment numbers, which are all irrelevant for your task. *You only need to look at four fields: **time stamp, source and destination IP addresses, and IP packet size**.* Note that some packets are not IP packets, meaning that the IP packet size field can be empty.

Step 1: Getting started

A basic structure of this assignment is given, you are expected to complete the **class design** for this assignment, but there are some fairly obvious parts to the design: Firstly, it's a GUI application with one window, so you'll need a JFrame. In all our GUI applications in the lectures, we have extended that JFrame by a dedicated subclass for the application. So you'll probably want to do the same here.

The Swing components are all associated with the JFrame, so you'll need to import the associated packages, configure the components and add them to the JFrame either directly or via their respective parent. The ListOWords and Album applications from the lectures/supporting material are good examples for how to do this. Don't forget to configure the ButtonGroup to link your radio buttons!

One exception is the JPanel that shows the data in tabular form. It would be best to extend JPanel, so a more specialised subclass can take care of the table drawing.

Start by placing all components visibly within the JFrame – you may wish to set the JPanel background colours to something a bit different so you can see where in the JFrame they end up being positioned. Remember that a larger vertical position value puts the component further *down* in the frame. Suggested values that work for me (but aren't compulsory):

- The JFrame has size 1000 by 500.
- I've put the JPanel with the radio buttons at position 0,0 (upper left corner of the JFrame's content pane) and made it 200 wide and 100 tall.
- The JPanel subclass with the table is at 0,100 (relative to the JPanel) and is 1000 wide by 1300 tall.

You can earn a total of **20 marks** in this step:

- JFrame subclass correctly started via a class implementing Runnable & using invokeLater(): **4 marks**
- Menu bar present, showing a File menu with the two required items: **4 marks.**
- Menu item "Quit" quits the application and Menu item "Open trace file" opens a JFileChooser: **4 marks**
- The two radio buttons are visible: **4 marks**

- The radio buttons are mutually exclusive and one is selected by default: **4 marks**

In order to get the marks for positioning, no components or labels must touch each other or overlap with other components (i.e., the JFrame and its contents must look neat and tidy).

Step 2: Extract the source and destination hosts from the trace file

First ensure that the combo box is invisible before you open the first trace file.

Using the trace file selected by the JFileChooser, parse the file to extract the source and destination hosts. Check which radio button is active and show the respective set of hosts selected in the combo box.

Hint: There are a few little catches here:

- 1) Splitting lines can be done with the `split()` method of the `String` class, which takes a regular expression as its parameter.
- 2) The lists must not contain duplicates, so you can't simply add each line's host entry to an `ArrayList` and then load that `ArrayList` into the `JComboBox`. Instead, Java has a data structure quite similar to an `ArrayList` that doesn't allow duplicates: the `HashSet` class. It implements the `Set` interface, and for `Strings`, you can use it as follows:

```
import java.util.Set;
import java.util.HashSet;
...
Set<String> myUniqueStrings = new HashSet<String>();
...
myUniqueStrings.add(aStringWeMayOrMayNotHaveSeenBefore);
```

If `aStringWeMayOrMayNotHaveSeenBefore` is not yet in the hash set `myUniqueStrings`, it will be added. If the string is already in the set, the `add()` method does nothing.

- 3) The lists must only contain IP addresses of IP version 4 (IPv4). Some packets in the trace file are **not** IP packets and the entries for source and destination in the respective lines will be empty. You must skip these packets. Think regular expressions, perhaps.
- 4) The lists in the combo box must be sorted. To do this, you need to figure out how to sort IP addresses in Java. This can be done in a number of ways. I use `Collections.sort()` and store the IP addresses in a class that has an appropriate `compareTo()` method.

Once you have accomplished this, ensure that the combo box updates every time you click the other radio button and every time you open a new trace file. For this, you need to:

- Implement and add the necessary event handlers for the radio button.
- Implement the code required to (where applicable) fetch the list of hosts (should you generate these lists once when you load the file, or each time the radio buttons change?).
- Implement the code required to wipe and re-populate the combo box. Look to the Album example from the lectures/supporting material for a guide here.

Ensure that you get the correct list of hosts in the combo box each time.

You can earn a total of **30 marks** in this step:

- The combo box is invisible before the trace file is selected and opened: **3 marks**
- The application shows some evidence of a trace file being opened and its contents being processed after one selects a file with the JFileChooser (e.g., table plot appearing, combo box getting populated with sensible data): **7 marks**
- The combo box becomes visible at this point in time: **4 marks**
- The combo box contains a list of IP addresses: **4 marks**

- The list contains the selected type of host addresses (192.168.0.x if source hosts are selected, 10.0.x.x for destination hosts): **4 marks**
- The list updates correctly when a radio button is chosen: **4 marks**
- The list updates correctly when we open a new trace file: **4 marks**

We use coderunner questions to test the correctness of data retrieved from the trace file. You can earn a total of **10 marks** in this step. Complete question 1 to question 5 in coderunner.

Step 3: Compute and show tabular data

Implement the necessary code to:

- Create a JTable and associated custom table model and add them to your custom JPanel. Note: you would also need to use a scroll pane (JScrollPane class) to accommodate all table records.
- Compute the total number of bytes (packet size) that the selected source/destination host sent/received. [*For the first of the two additional table rows*]
- Compute the average number of bytes across all data rows. [*For the second additional row*]
- In order to fulfil the requirements of editable cells and other cell rendering requirements, you would need to create your own custom table model (model in MVC), which can be performed by creating a subclass of AbstractTableModel class of the Java API. You should then be able to link this custom table model to your JTable (view in MVC) by setting your custom model object as the table model for your JTable object.
- Use table cell rendering to improve the aesthetics of the table view.
- Ensure that the changes made to the editable fields are visible in the trace files, whenever user presses the 'enter' key (or 'return' key) after editing a cell value.
- Useful links:
 - <https://docs.oracle.com/javase/tutorial/uiswing/components/table.html>
 - <https://docs.oracle.com/javase/7/docs/api/javax/swing/table/TableModel.html>
 - <https://docs.oracle.com/javase/tutorial/displayCode.html?code=https://docs.oracle.com/javase/tutorial/uiswing/examples/components/TableDemoProject/src/components/TableDemo.java>
 - http://www2.hawaii.edu/~takebaya/ics111/jtable_custom/jtable_custom.html

You can earn a total of **45 marks** in this step:

- The application shows a table of sorts: **3 marks**
- The data table is actually correct (as per the text file) and corresponds to the host selected: **3 marks**
- The table updates correctly when a radio button is chosen: **2 marks**
- The IP address column name updates correctly when a radio button is chosen: **2 marks**
- The table updates correctly when an IP address is chosen: **2 marks**
- The table updates correctly when we open a new trace file: **3 marks**
- The table shows all relevant column and row values (i.e. table rows and columns are scrollable): **3 marks**
- There are two additional rows at the end of the table representing the Average and Total IP Packet Size for each table view: **6 marks**
- The two additional rows should be rendered in a different color than the rest of the table: **3 marks**
- Innovative look and feel of the table structure and table data: **3 marks**
- The table cell values of the last column (IP packet size) are editable: **3 marks**
- The changes made to the editable fields are visible in an external text file (i.e. user saves all valid packets into a text file): **12 marks**

We use coderunner questions to test the correctness of the data in the table. You can earn a total of **10 marks** in this step. Complete question **6** to question **7** in coderunner.

Step 4: Document your classes

Document your design. The documentation must consist of:

- 1) "Javadoc" style comments for each public non-inherited method in your code (except for event handlers in anonymous classes). See "Writing Doc Comments" under: <http://www.oracle.com/technetwork/java/javase/tech/index-137868.html>
Your comments only need a description of what the method does and the @param and @return block tags.

You can earn a total of **5 marks** in this step:

- *Javadoc comments*: **5 marks** (minus one mark for each un-/incorrectly documented method)

Debugging

I strongly recommend that you use Eclipse for this project. One particular challenge is the parsing of the trace file. Much of the intermediate operations here happen behind the scenes – there is no visible output till the very end of a lot of these operations and many things that can potentially go wrong. One nice feature in Eclipse is that you can still use `System.out.println()` etc. to write to the console, so you can add such statements to check that your code behaves as expected. You'll also get to see any stack traces when exceptions are thrown. And of course, you can add breakpoints and debug using Eclipse's debugging features!

ACADEMIC INTEGRITY

The purpose of this assignment is to help you develop a working understanding of some of the concepts you are taught in the lectures.

We expect that you will want to use this opportunity to be able to answer the corresponding questions in the tests and exam.

We expect that the work done on this assignment will be your own work.

We expect that you will think carefully about any problems you come across, and try to solve them yourself before you ask anyone for help.

This really shouldn't be necessary, but **note the comments below about the academic integrity related to this assignment:**

The following sources of help are acceptable:

- Lecture notes, tutorial notes, skeleton code, and advice given by us in person or online, with the exception of sample solutions from past semesters.
- The textbook.
- The official Java documentation and other online sources, as long as these describe the general use of the methods and techniques required, and do not describe solutions specifically for this assignment.
- Piazza posts by, or endorsed by an instructor.
- Fellow students pointing out the cause of errors in your code, without providing an explicit solution. The following sources of help are not acceptable:
- Getting another student, friend, or other third party to instruct you on how to design classes or have them write code for you.
- Taking or obtaining an electronic copy of someone else's work, or part thereof.
- Give a copy of your work, or part thereof, to someone else.
- Using code from past sample solutions or from online sources dedicated to this assignment.

The Computer Science department uses copy detection tools on all submissions. Submissions found to share code with those of other people will be detected and disciplinary action will be taken. To ensure that you are not unfairly accused of cheating:

- Always do individual assignments by yourself.
- Never give any other person your code or sample solutions in your possession.
- Never put your code in a public place (e.g., Piazza, forum, your web site).
- Never leave your computer unattended. You are responsible for the security of your account.
- Ensure you always remove your USB flash drive from the computer before you log off, and keep it safe.

Concluding Notes

Not everything that you will need to complete this assignment has or will be taught in lectures of the course. You are expected to use the Oracle Java API documentation and tutorials as part of this assignment.

Post any questions to Piazza or take them along to the COMPSCI 230 labs/lectures – this way, the largest number of people get the benefit of insight!

I may comment along the lines on Piazza to deal with any widespread issues that may crop up.

As an additional challenge: As long as the functionality / specifications above are preserved, you are welcome to add any additional features that may be useful, e.g.,

- The ability to look at data packet flow between a particular source and destination host
- The ability to look at data packet flow from/to/between particular ports
- The ability to print the table / data in the table, etc.

Any such additional features do not attract marks but would certainly count in your favour should you find yourself in the unfortunate position of needing an exam aegrotat, or the fortunate position of needing a recommendation letter from me.